



# CSL2050

# Pattern Recognition & Machine Learning

## Course Project

# “Face Identification”



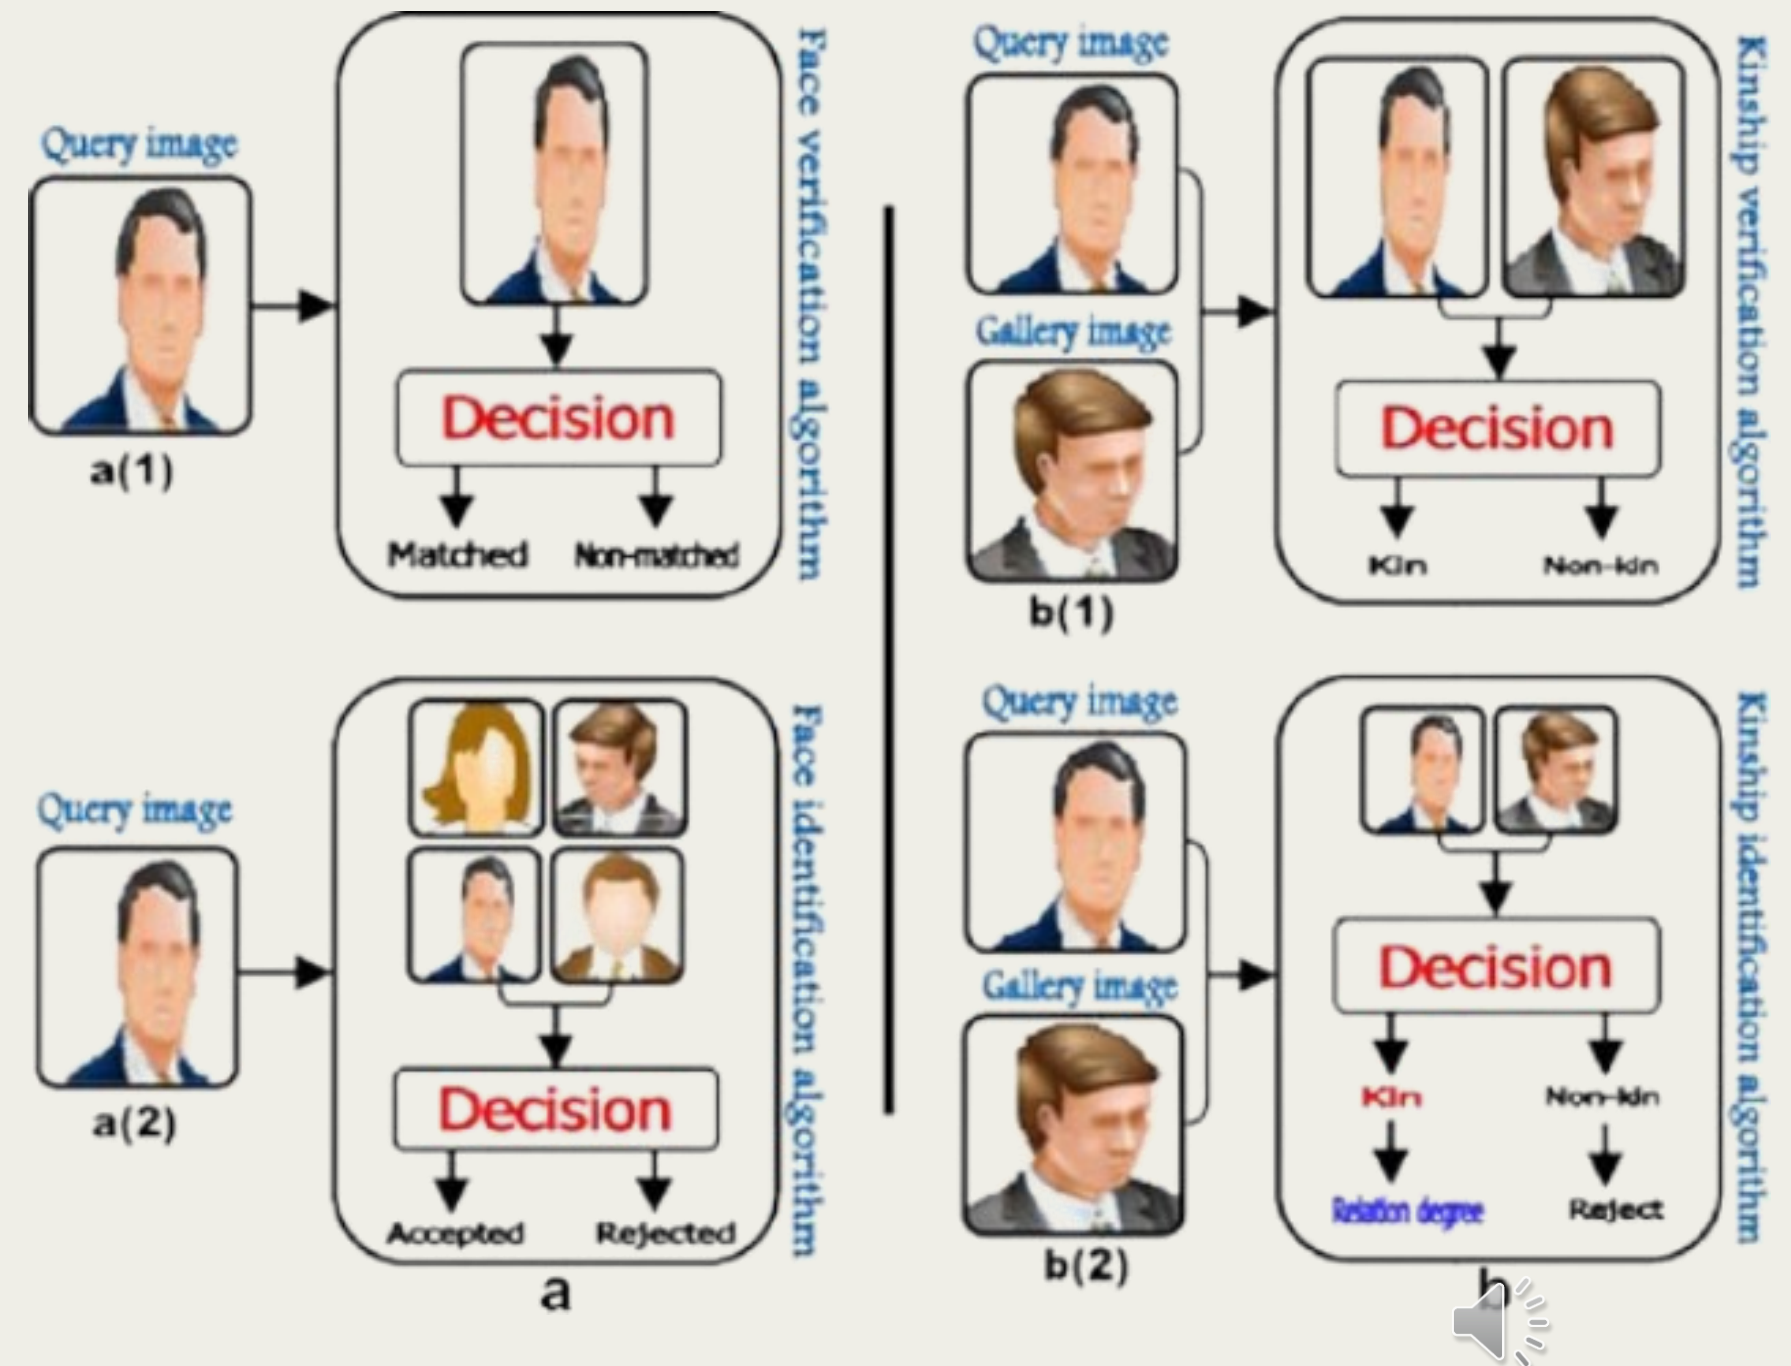
Sandeep Soni, Atanu Kayal, Japneet Singh,  
Varchasva Saxena, Vishal



# Introduction & Motivation

## Problem Statement:

Build a robust face-identification system that classifies any given face image into one of K known identities



# Introduction & Motivation

**Problem Statement:** Build a robust face-identification system that classifies any given face image into one of  $K$  known identities

## Key Application:

**Automated Attendance – mark people in/out in real time using face recognition**



# Datasets

## LFW Dataset

(Labeled Faces in the Wild)

- **~13,000 images of 5,749 individuals**
- **Captured in uncontrolled settings**

### Problems:

- **Few images per identity → overfitting**
- **Imbalanced classes → biased models**
- **Poor generalization for deep models**

**Only good for face verification, not identification**  
**Initially used for prototyping**

## VGGFace2

(Final Dataset)

- **3.3M images across 9,131 identities**
- **High intra-class variation**
- **Balanced distribution**

### Our subset:

- **50 identities**
- **300–350 images each (total ≈16,000 images)**
- **Cleaned and preprocessed for consistency**





# Datasets

## LFW Dataset

(Labeled Faces in the Wild)

- ~13,000 images of 5,749 individuals
- Captured in uncontrolled settings

### Problems:

- Few images per identity → overfitting
- Imbalanced classes → biased models
- Poor generalization for deep models

Only good for face verification, not identification  
Initially used for prototyping

## VGGFace2

(Final Dataset)


- **3.3M images across 9,131 identities**
- **High intra-class variation**
- **Balanced distribution**

### Our subset:

- **50 identities**
- **300–350 images each (total ≈16,000 images)**
- **Cleaned and preprocessed for consistency**




# Preprocessing Pipeline




**Original Image**




# Preprocessing Pipeline




**Detected bounding  
box overlay**



# Preprocessing Pipeline



**Cropped & resized  
face Image**



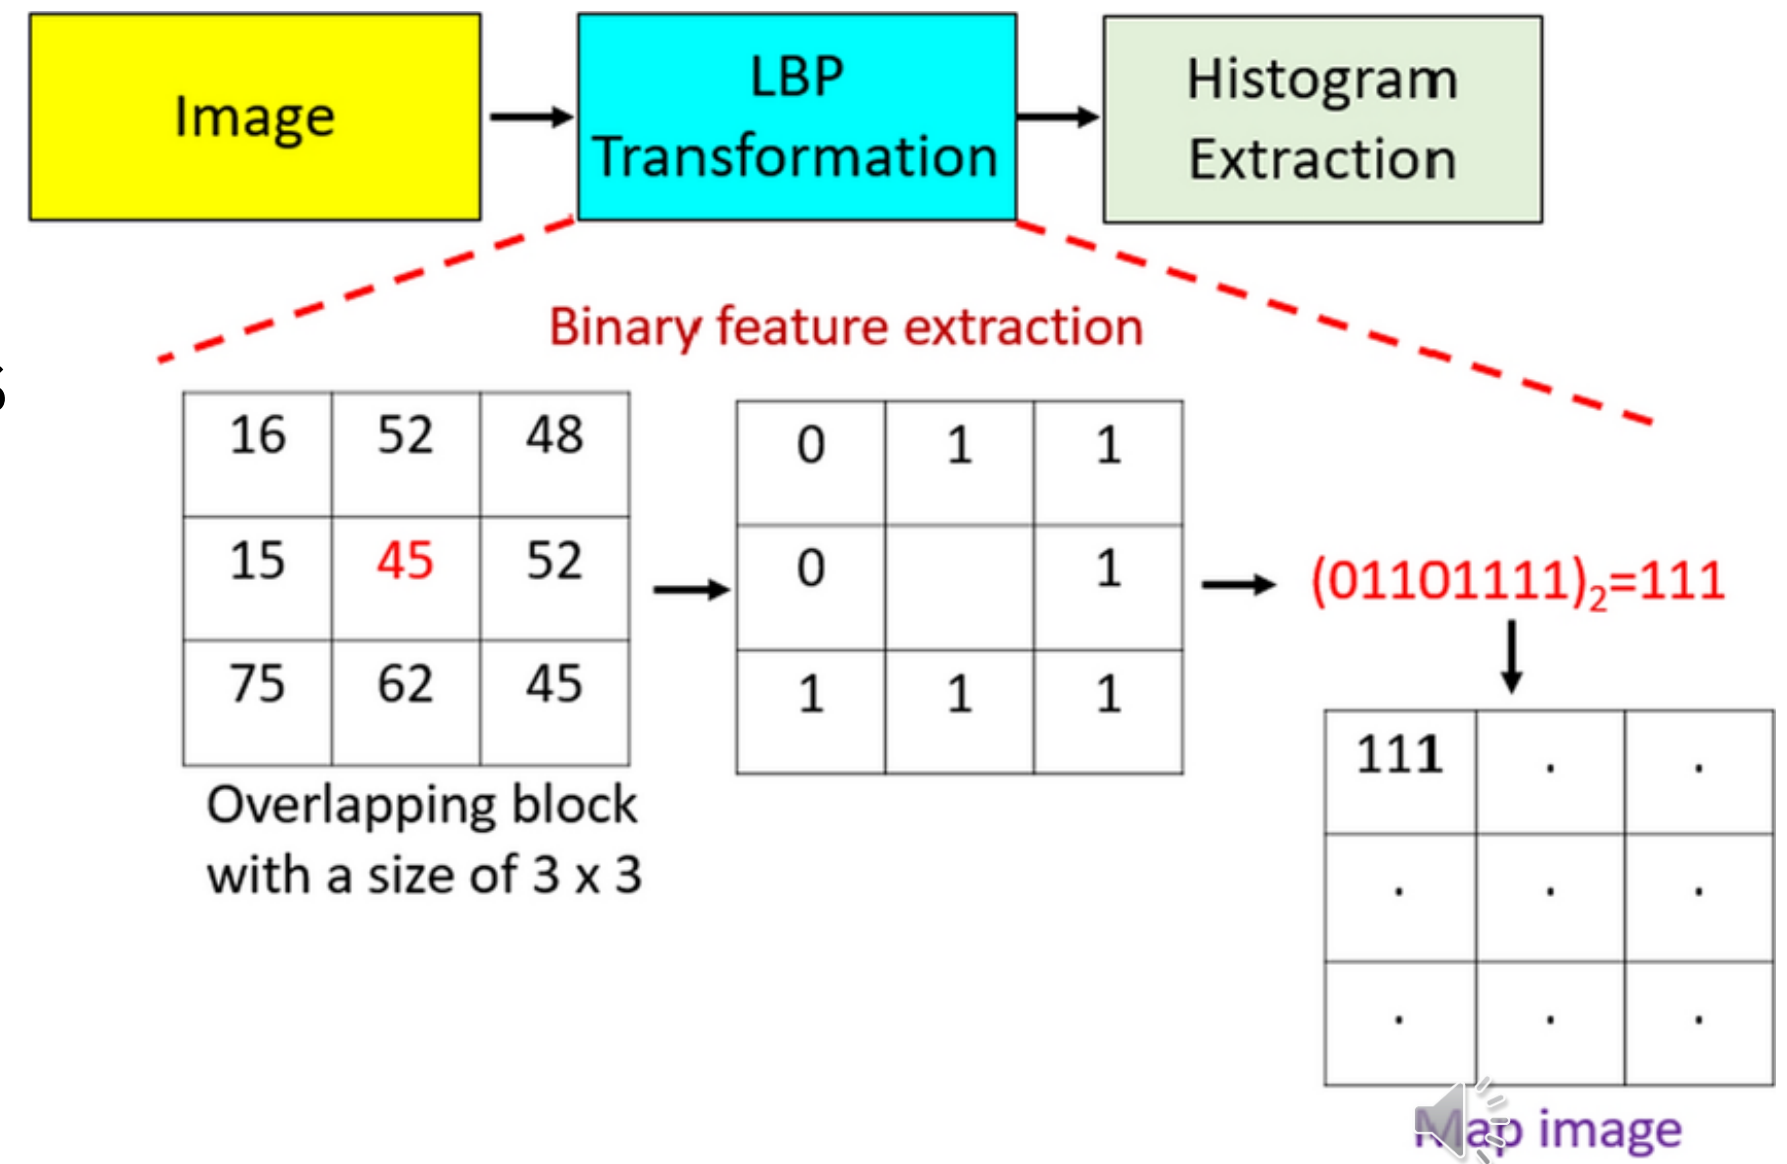


# Feature Extraction Techniques

Convert face images into numerical feature vectors that models can learn from. Each method captures different aspects of facial information.

## 1 LBP (Local Binary Patterns)

- Captures local texture patterns
- Robust to lighting variation
- Outputs: Histogram of binary patterns
- Works on grayscale images
- Fast but limited context

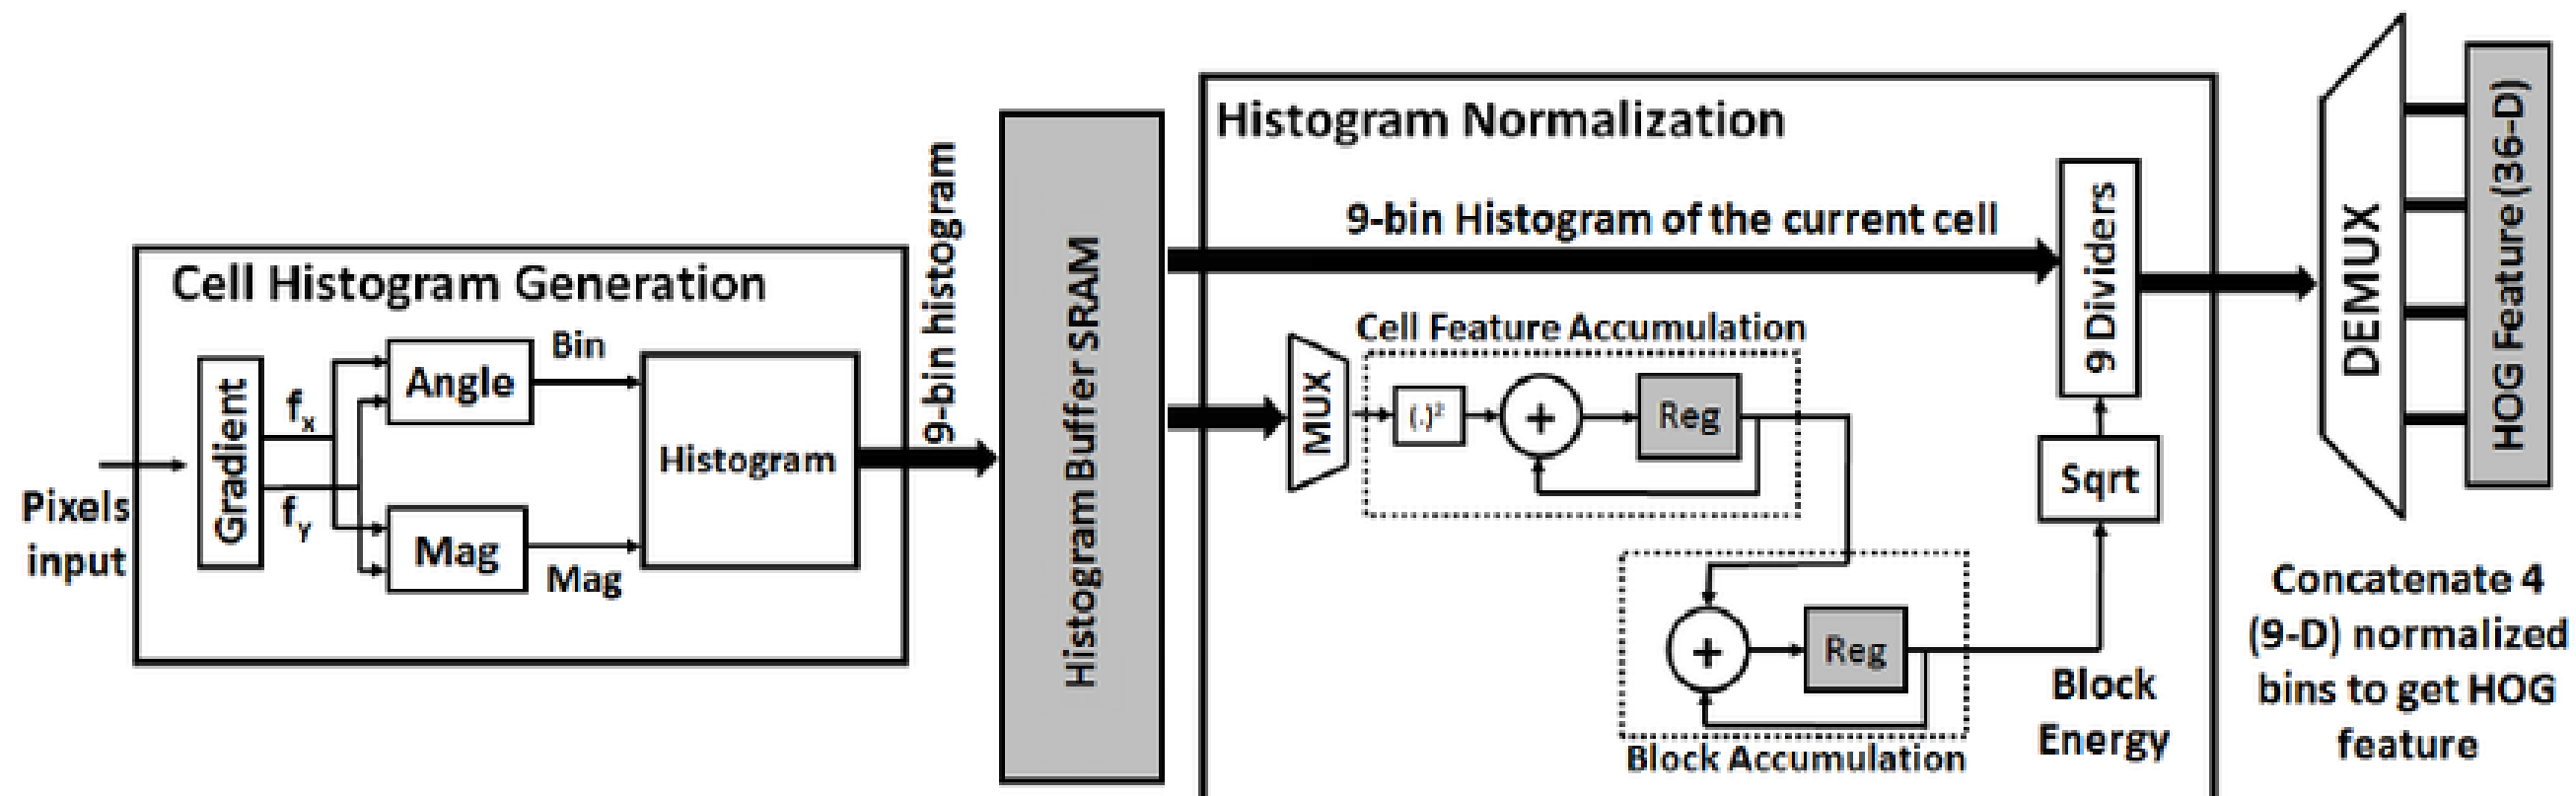


# Feature Extraction Techniques

Convert face images into numerical feature vectors that models can learn from. Each method captures different aspects of facial information.

## 2 HoG (Histogram of Oriented Gradients)

- Captures edge orientation & structure
- Divides image into cells, computes gradient histograms
- Good for shape information
- Lacks semantic detail



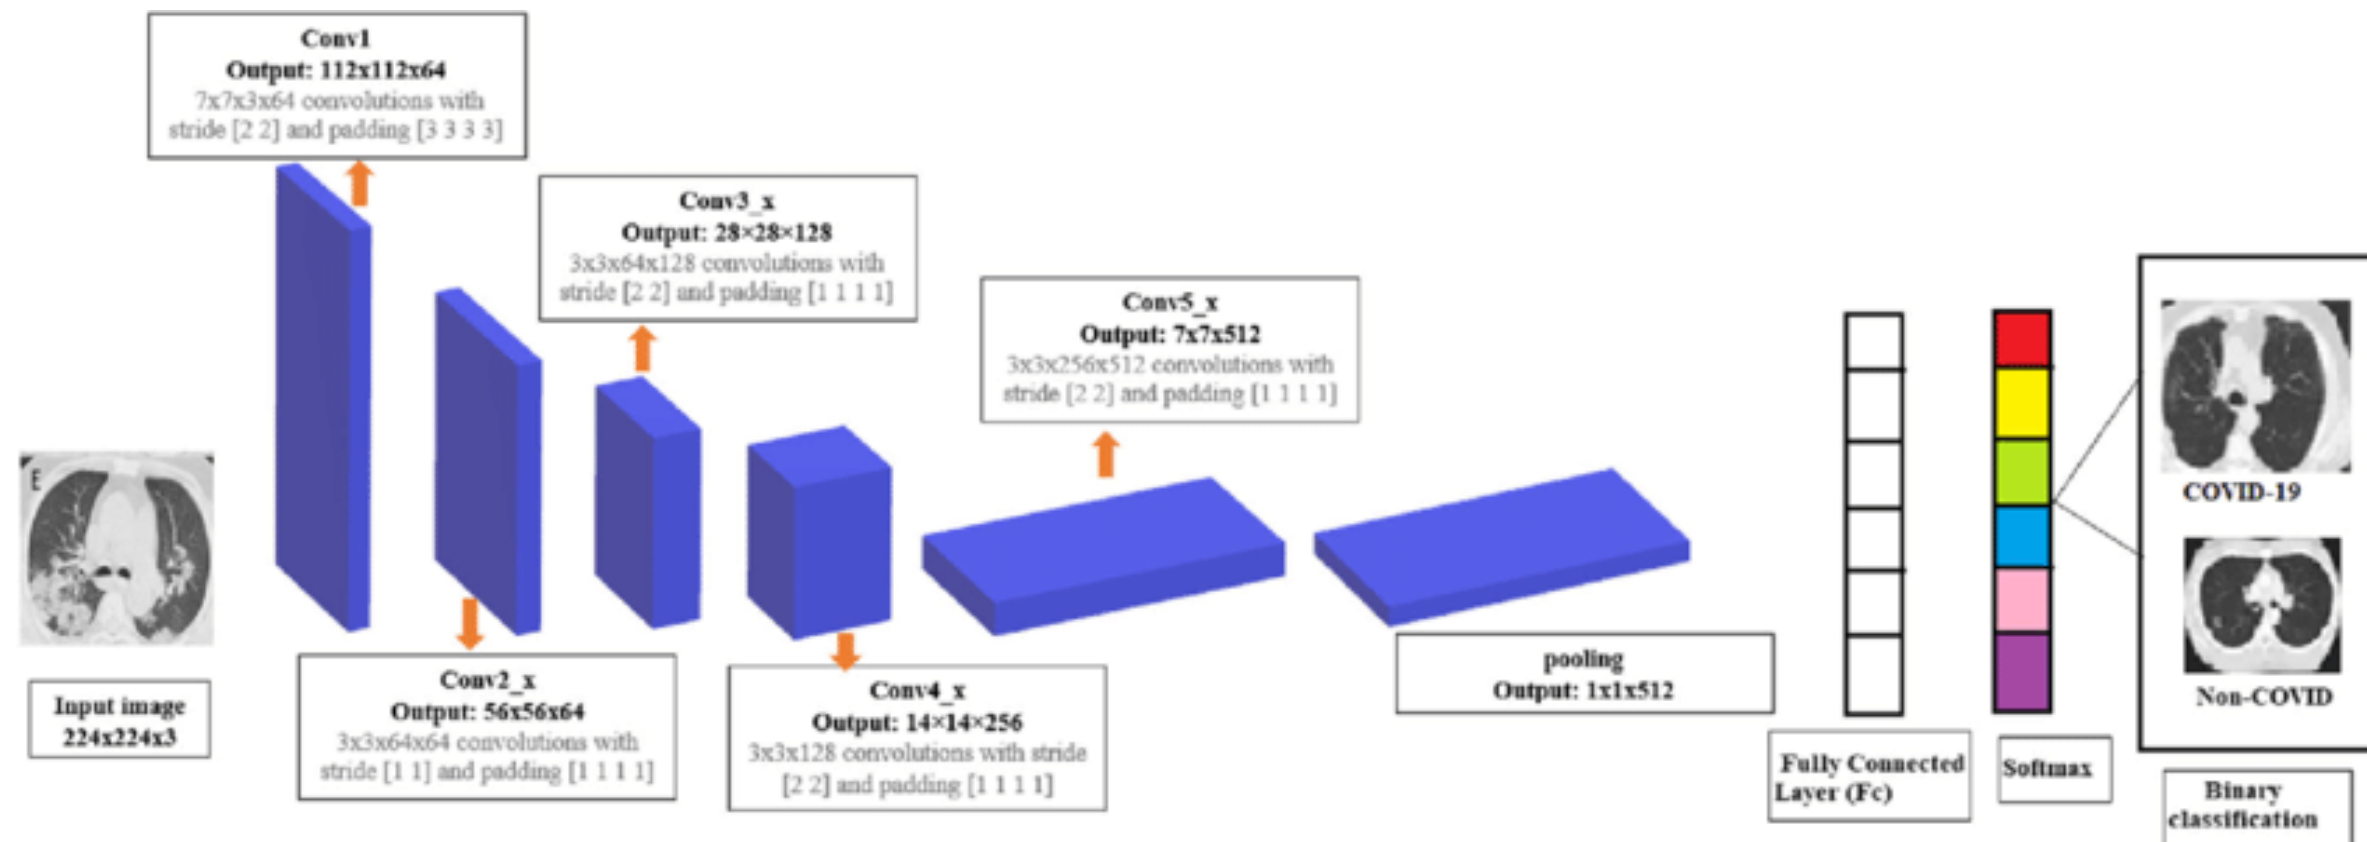
# Feature Extraction Techniques

Convert face images into numerical feature vectors that models can learn from. Each method captures different aspects of facial information.

3

## CNN Features (ResNet18)

- Deep semantic feature embeddings
- Pretrained on large datasets
- Outputs: 512-dimensional vector
- Most accurate & robust , uses RGB images






# Feature Extraction Techniques

Technique	Sample Image	Output
LBP	Cropped face	LBP pattern histogram (grayscale overlay)
HoG	Same face	Gradient cell overlay visualization
CNN	Same face	Activation heatmap / embedding visualization (e.g., PCA projected)

CNN Feature Maps



Original Image




LBP Image



Original Image



HoG Image



# Feature Fusion

**Concatenation: LBP + HoG + CNN vector**

**Why ?**

- LBP captures textures
  - HoG adds shape details
  - CNN provides deep semantics
- Fusion creates a rich representation of each face






# Dimensionality Reduction Techniques:


## PCA (Principal Component Analysis)

- Projects data to lower dimensions
- Keeps variance
- Improves training time

PCA: component axes that maximize the variance



LDA: maximizing the component axes for class-separation



## LDA (Linear Discriminant Analysis)

- Supervised reduction
- Maximizes class separation
- Better for classification



# Classification Techniques

## How We Classified the Faces ?

Once we had our feature vectors (from LBP, HoG, CNN, or combined), we experimented with various machine learning and deep learning classifiers to predict the correct identity.

### Traditional ML Models (using features):

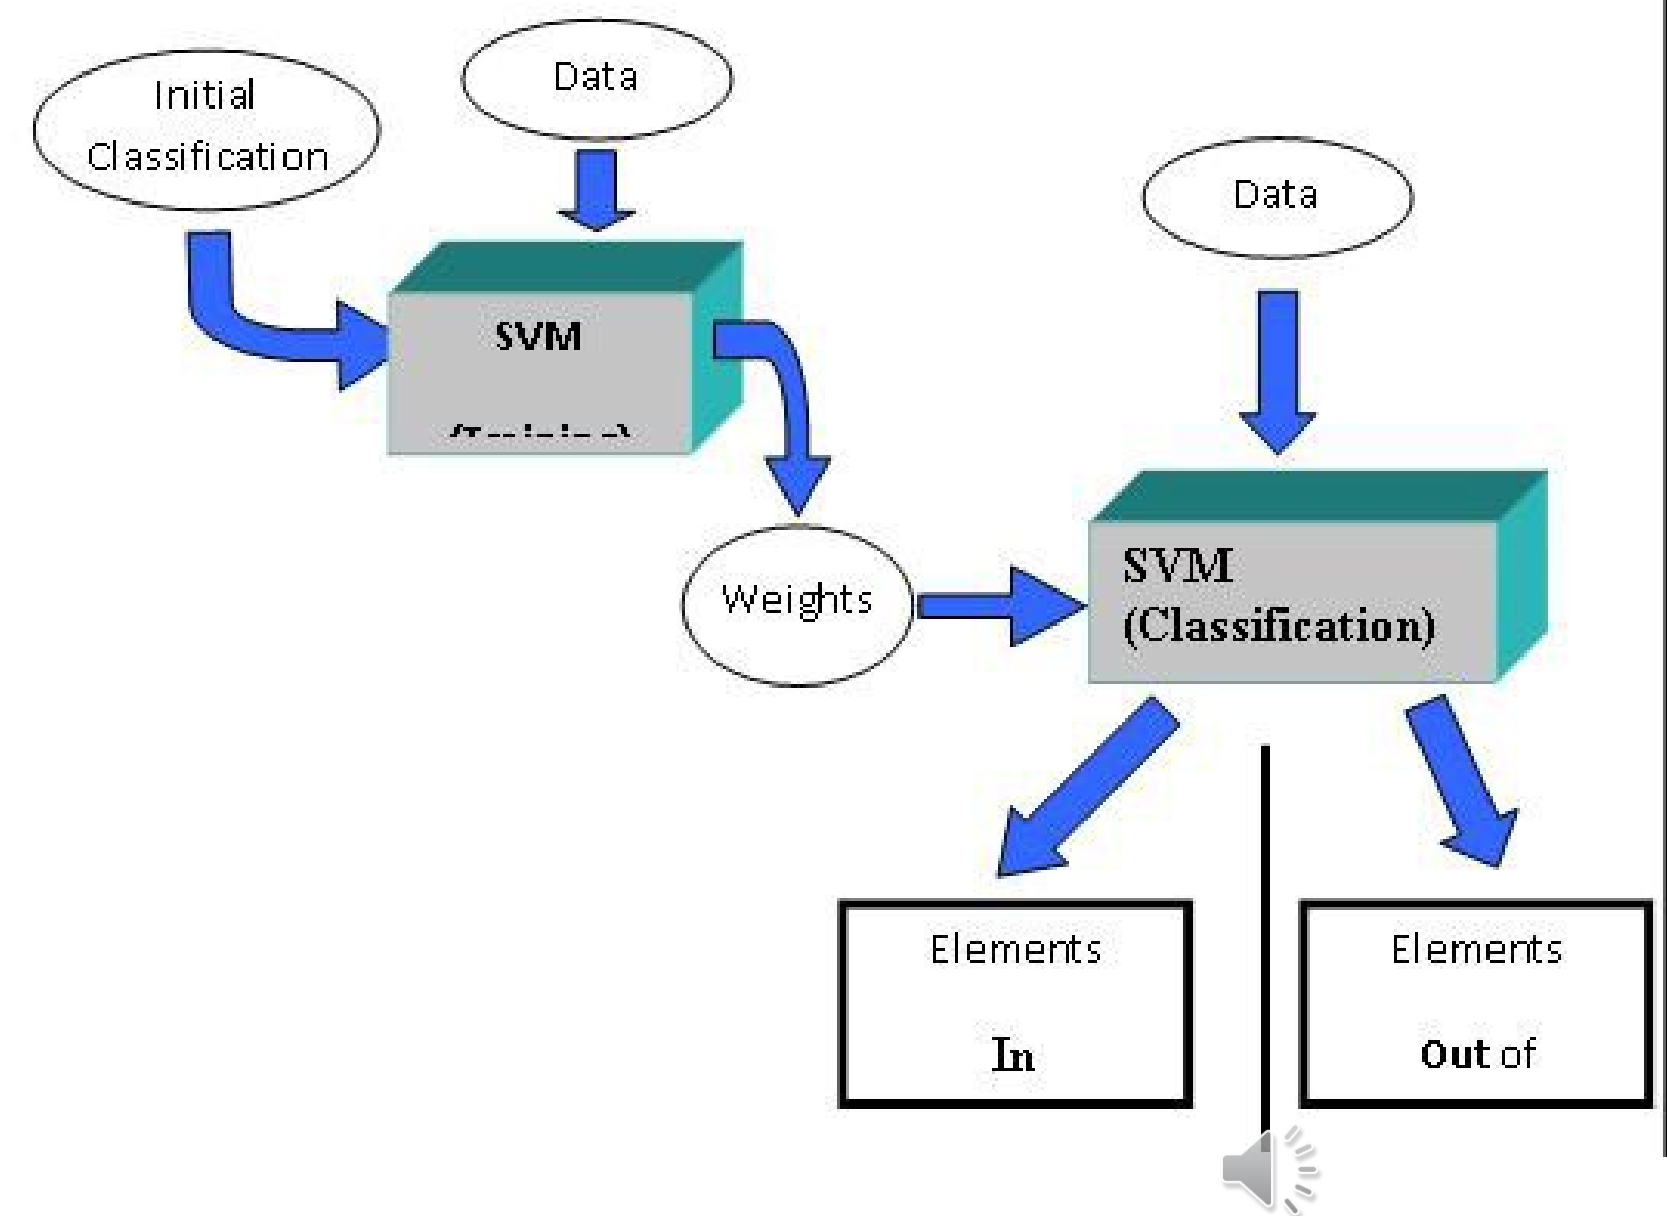
- SVM (Support Vector Machine)
- KNN (K-Nearest Neighbors)
- Logistic Regression
- Random Forest & XGBoost
- ANN (Shallow Neural Network)



# Classification Techniques

## 1. SVM (Support Vector Machine)


- Best performer on extracted features
- Handles high-dimensional data well
- Kernel used: RBF
- Accuracy: 76%



# Classification Techniques

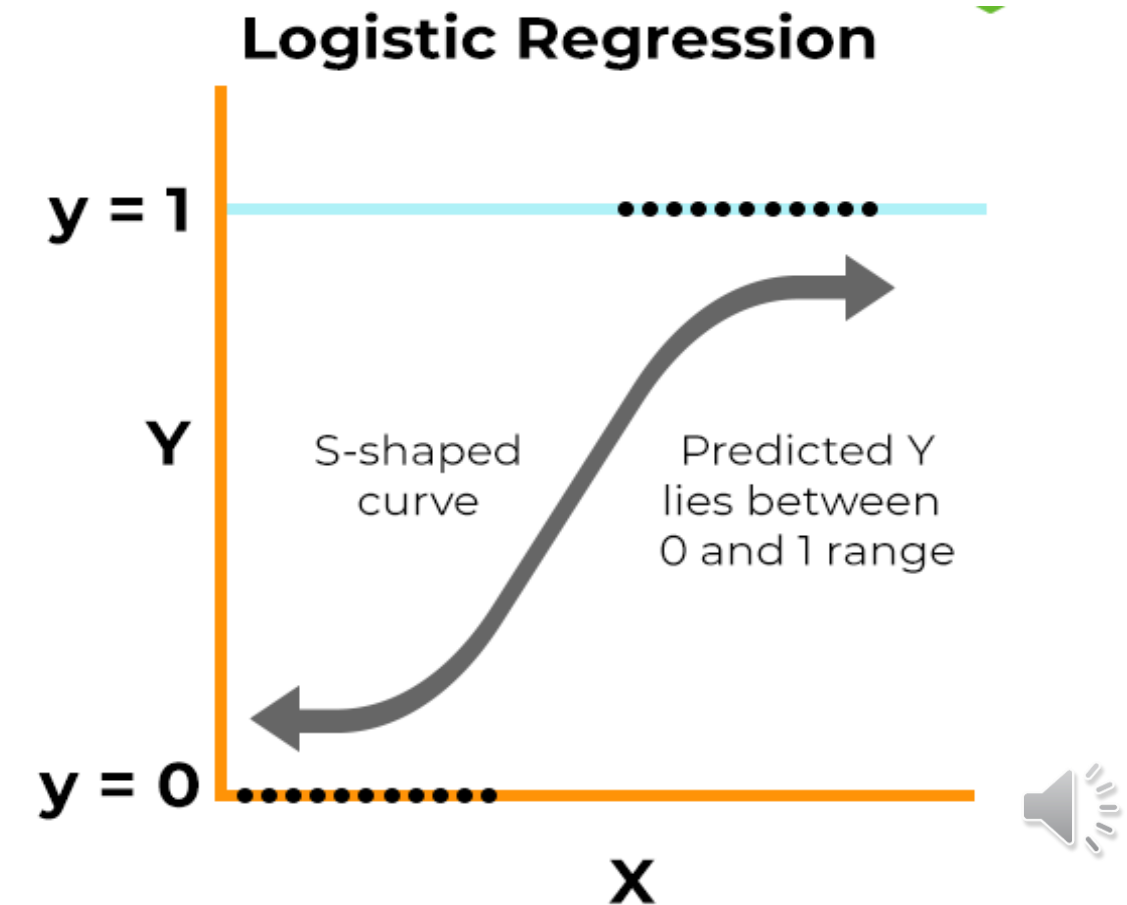
## 2. KNN (K-Nearest Neighbors)

- Distance-based, simple
- Accuracy: 58%



## 3. Logistic Regression

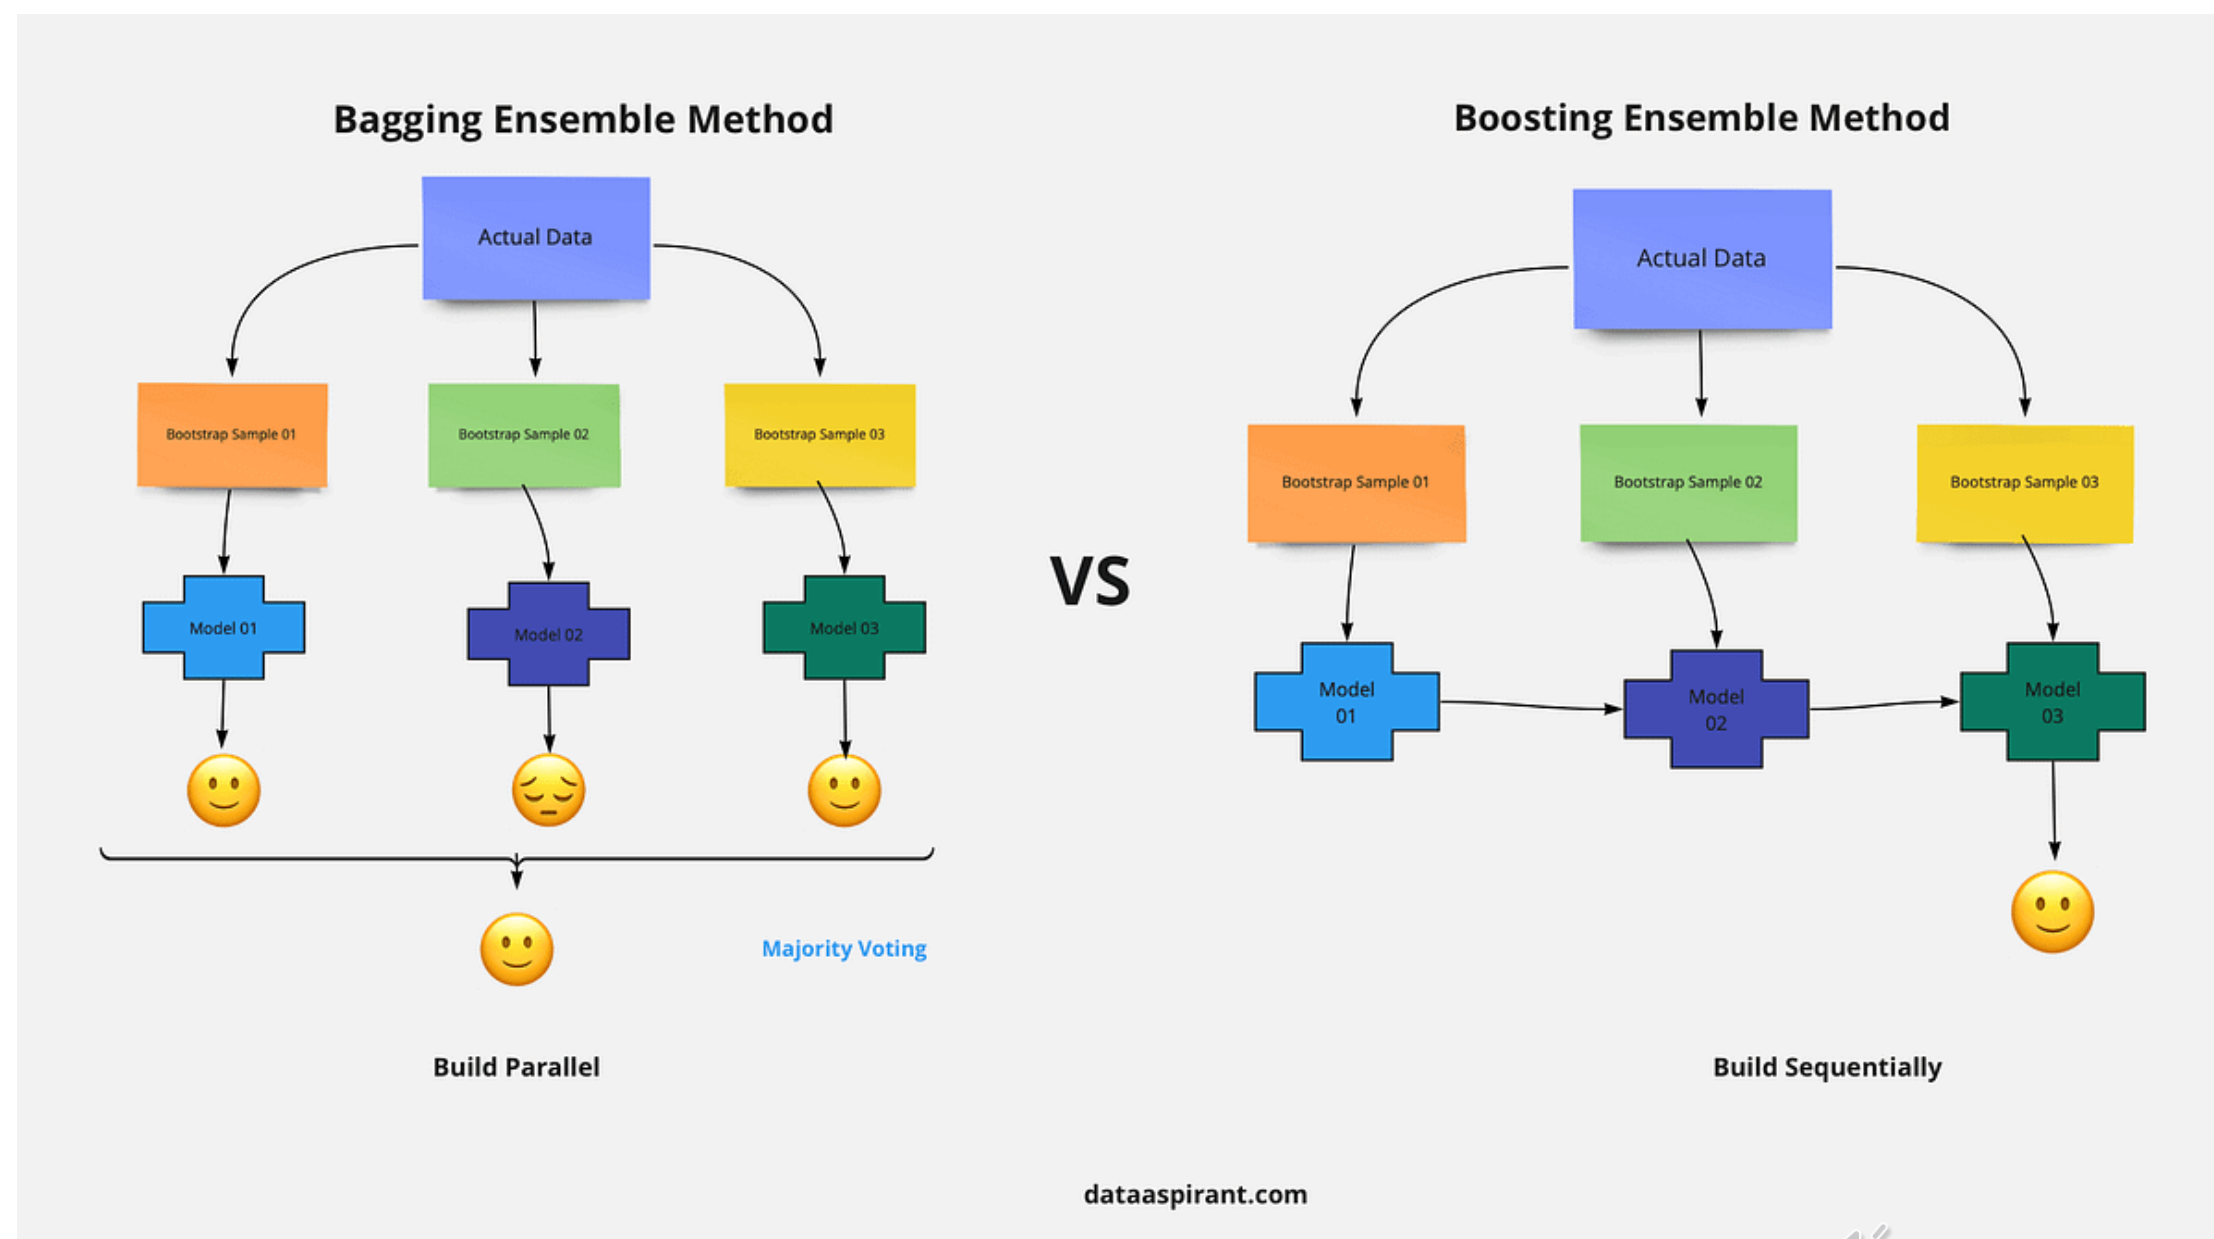
- Probabilistic model
- Accuracy: 66%



# Classification Techniques

## 4. Random Forest & XGBoost :

- Ensemble-based
- Struggled with high-dimensional fused features
- RF Accuracy: 57%, XGBoost: 63%

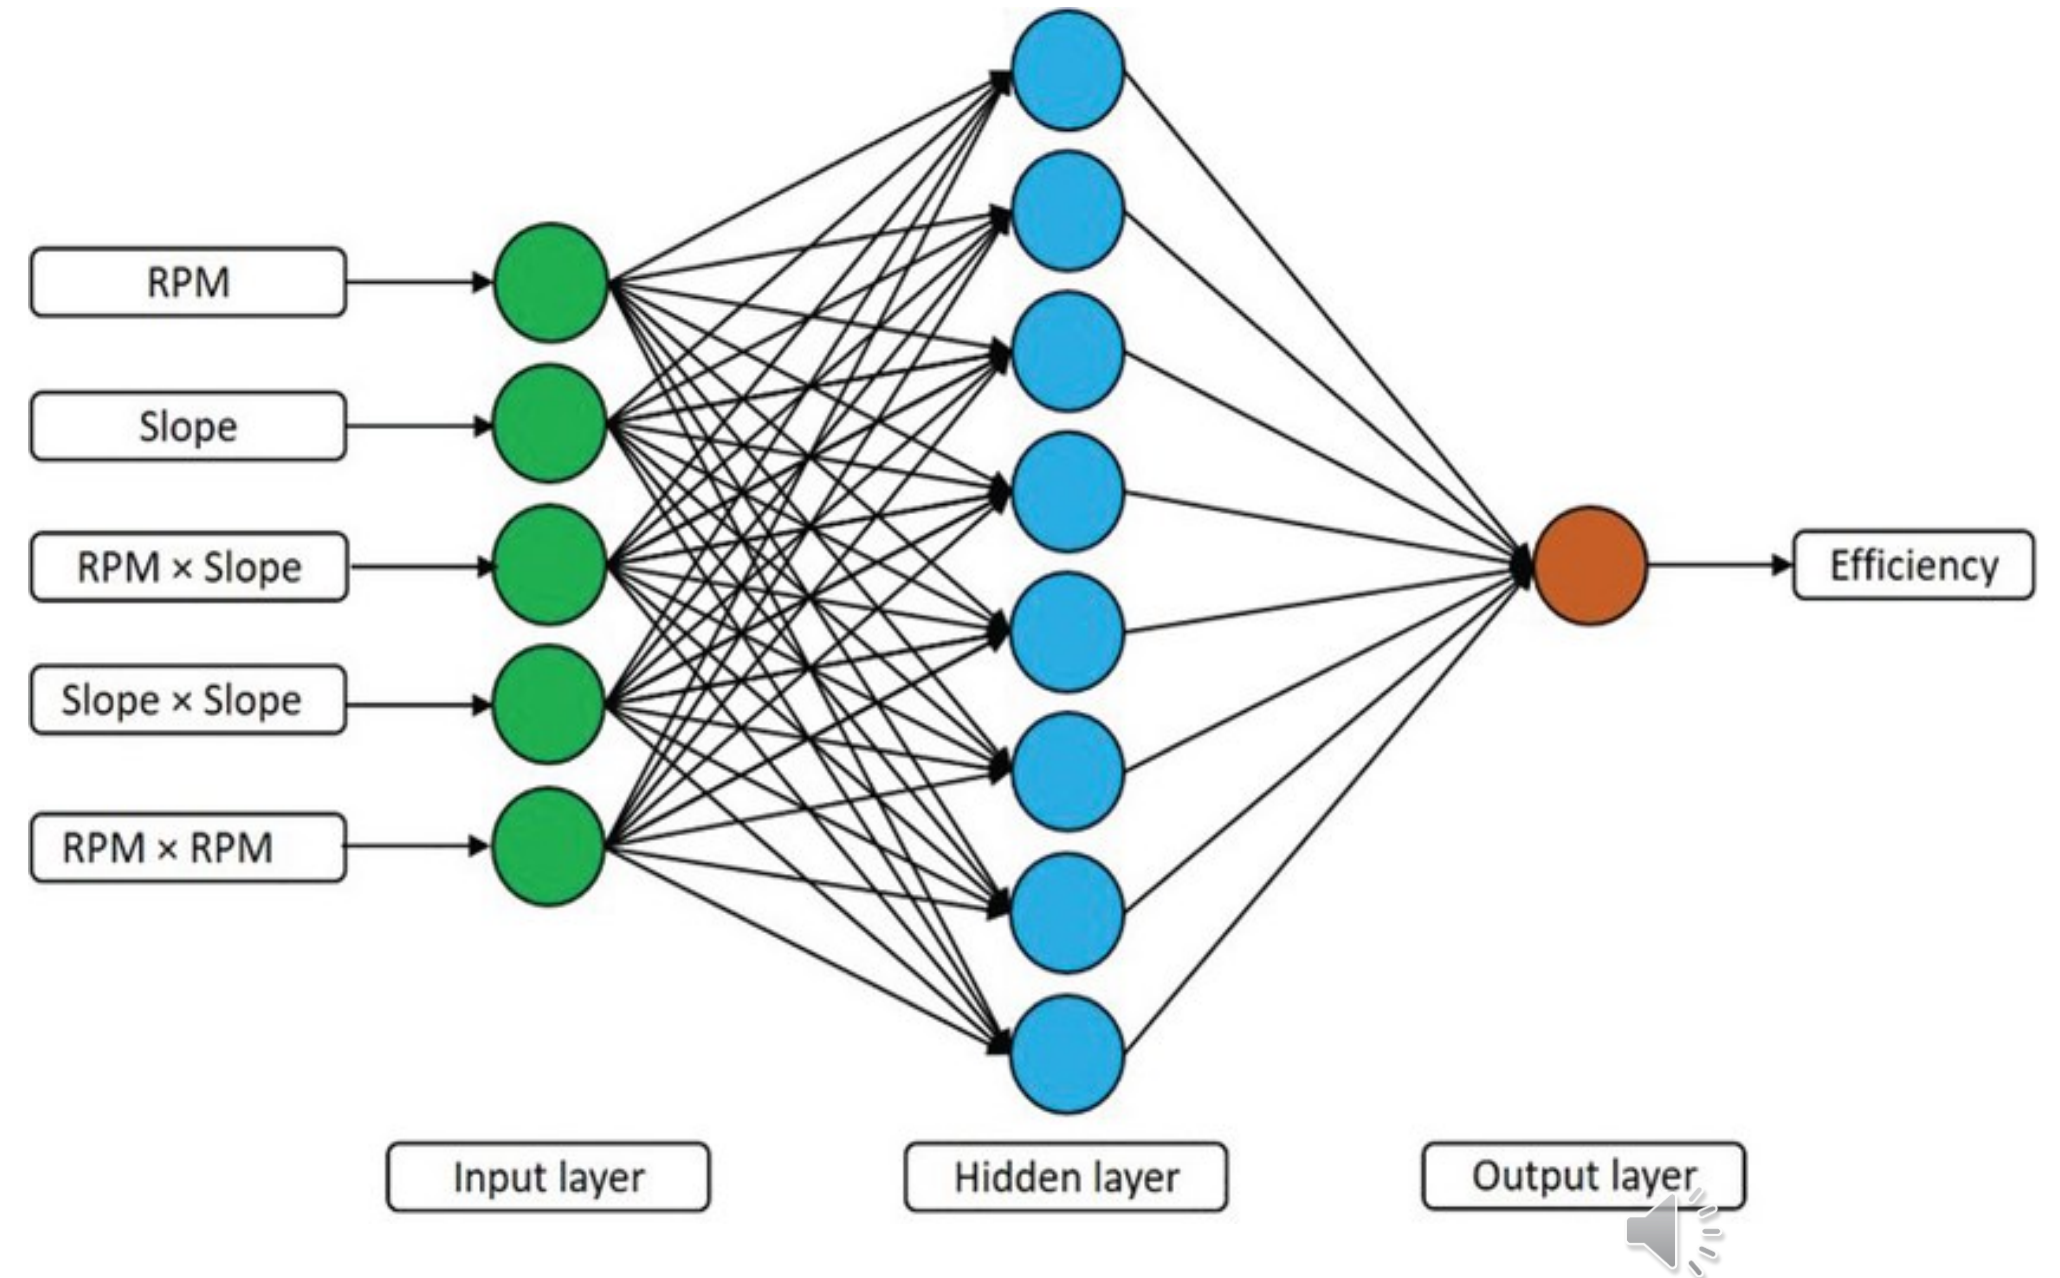





# Classification Techniques

## 5. ANN (Artificial Neural Network)


- 2 hidden layers with ReLU
- Accuracy: 63%



# Accuracy with Different Features (using SVM):



# Accuracy with Different ML Models:



# Best Approach & Trade-Offs

## Actual Best Model Implemented:

*Support Vector Machine (SVM)* – Achieved **76% accuracy**, outperforming other classical ML models.

## Top performer Model: ResNet50 (92% Accuracy)

*ResNet50* showed the highest simulated accuracy at 92%, but was not implemented in our project.


## Tradeoffs :

### SVM (Used):

- ✓ Simple, fast, interpretable
- ✗ Lower accuracy on complex data

### ResNet50 (not used ):

- ✓ Highest accuracy (92%), learns features automatically
- ✗ Complex, needs more data & compute



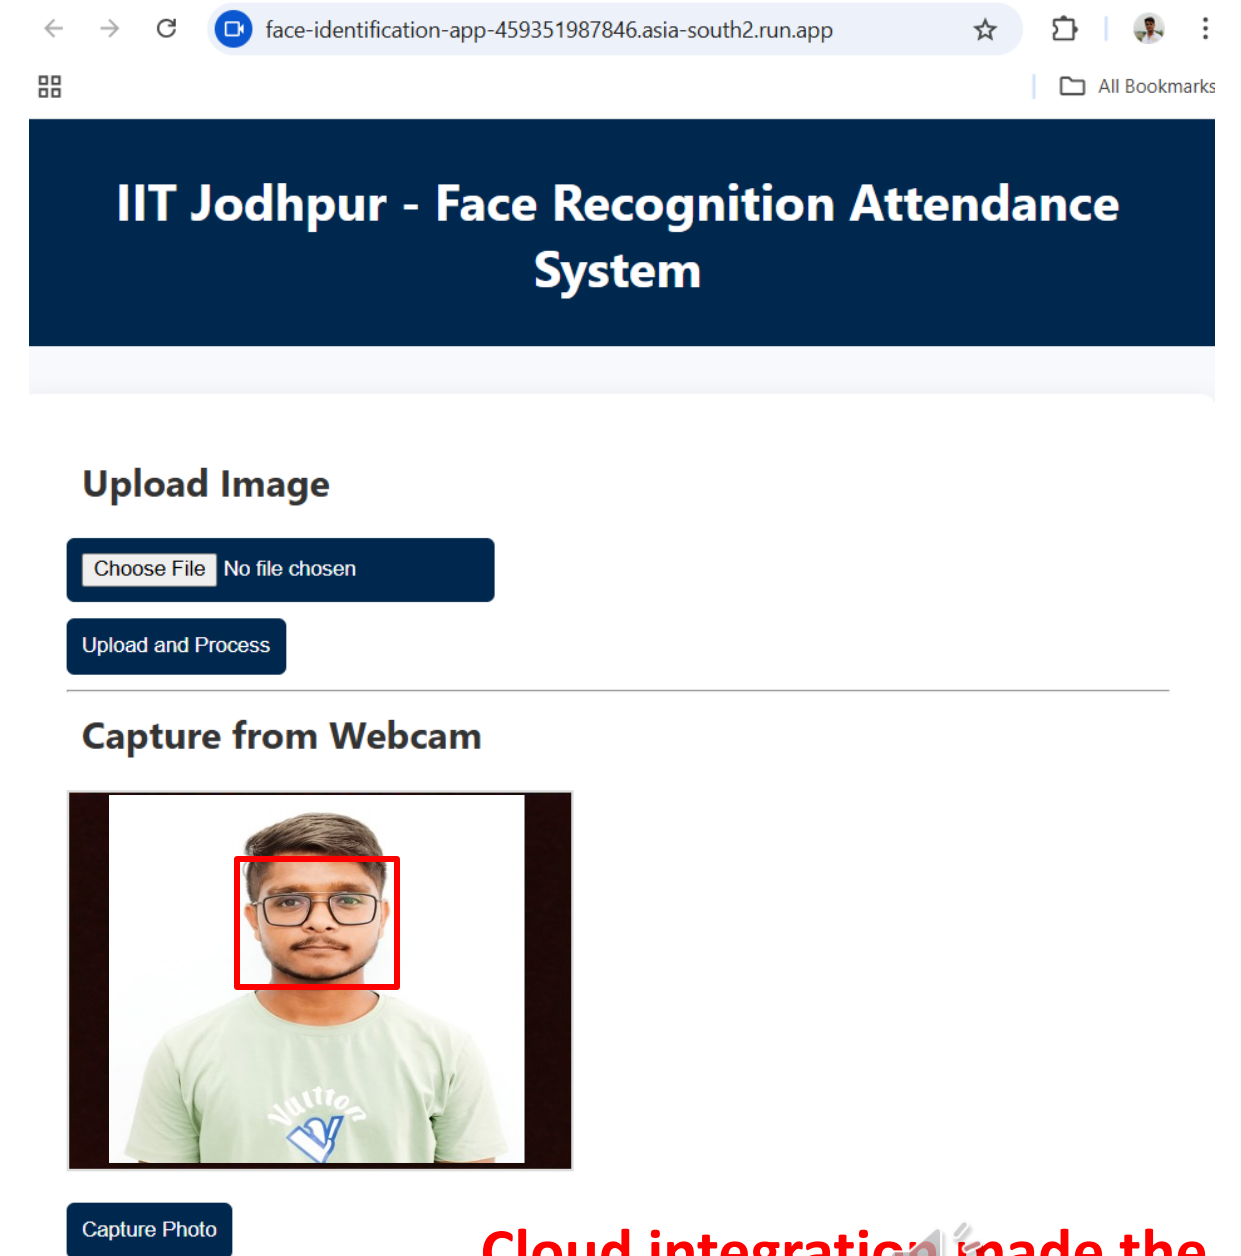
# Final Summary & Deployment

## System Overview:

- Real-time face recognition attendance using SVM
- Features extracted via LBP, HoG, and simulated CNN
- GUI for easy user interaction and attendance marking

## Google Cloud Deployment:

- **Firebase:** Real-time attendance updates & secure database
- **Cloud Storage:** Stores models & face data for fast access
- **Web Hosting:** Deployed project interface on GCP
- **Scalable & remotely accessible**



**Cloud integration made the system scalable, fast, and accessible from anywhere.**



😊 **Thankyou !** 😊

